

LoRA inside a 4-bit QLoRA layer

Frozen quantized backbone · trainable low-rank update

$$h = \text{dequant}(W_0^{\text{NF4}})x + \frac{\alpha}{r}BAx$$

